



Hierarchical spatial modelling for applied population and community ecology

Jeffrey W. Doser, Marc Kéry,
Gesa von Hirschheydt

24-27 June 2024





Modelling big spatial data

Jeffrey W. Doser
24-27 June 2024

Bayesian spatial linear model

$$y(\mathbf{s}_j) \sim \text{Normal}(\mu(\mathbf{s}_j), \tau^2)$$

$$\mu(\mathbf{s}_j) = \beta_0 + \beta_1 \cdot x_1(\mathbf{s}_j) + \mathbf{w}(\mathbf{s}_j)$$

$$\mathbf{w} \sim \text{Multivariate Normal}(\mathbf{0}, \mathbf{C}(d, \sigma^2, \phi))$$

- Recall our use of Gaussian Processes (GPs) to model the spatial random effects
- Values of \mathbf{w} are determined by the covariance matrix \mathbf{C} .
- \mathbf{C} is a $J \times J$ spatial covariance matrix (J is the number of spatial locations)

Covariance matrix with $J = 5$

$$\begin{bmatrix} k_{1,1} & k_{1,2} & k_{1,3} & k_{1,4} & k_{1,5} \\ k_{2,1} & k_{2,2} & k_{2,3} & k_{2,4} & k_{2,5} \\ k_{3,1} & k_{3,2} & k_{3,3} & k_{3,4} & k_{3,5} \\ k_{4,1} & k_{4,2} & k_{4,3} & k_{4,4} & k_{4,5} \\ k_{5,1} & k_{5,2} & k_{5,3} & k_{5,4} & k_{5,5} \end{bmatrix}$$



Covariance
between site 5 and
site 3

Covariance matrix with $J = 5$

$$\begin{bmatrix} k_{1,1} & k_{1,2} & k_{1,3} & k_{1,4} & k_{1,5} \\ k_{2,1} & k_{2,2} & k_{2,3} & k_{2,4} & k_{2,5} \\ k_{3,1} & k_{3,2} & k_{3,3} & k_{3,4} & k_{3,5} \\ k_{4,1} & k_{4,2} & k_{4,3} & k_{4,4} & k_{4,5} \\ k_{5,1} & k_{5,2} & k_{5,3} & k_{5,4} & k_{5,5} \end{bmatrix}$$

What happens when
 $J = 200$? $1,000$?
 $100,000$?

Covariance
between site 5 and
site 3



Gaussian process

- Flexible approach to account for spatial autocorrelation

Gaussian process

- Flexible approach to account for spatial autocorrelation
- But... becomes extremely slow as the number of sites increases
- Not practical for data sets with hundreds of data points, let alone thousands.
- Computational bottleneck: dealing with a large, *dense* $J \times J$ matrix
- The "big n" problem



The "big n" problem underneath the hood



The "big n" problem underneath the hood

- Estimating the parameters in our model requires calculating:

$$p(\mathbf{w}) \propto -\frac{1}{2}\log(\det(\mathbf{C})) - \frac{1}{2}\mathbf{w}^\top \mathbf{C}^{-1}\mathbf{w}$$



The "big n" problem underneath the hood

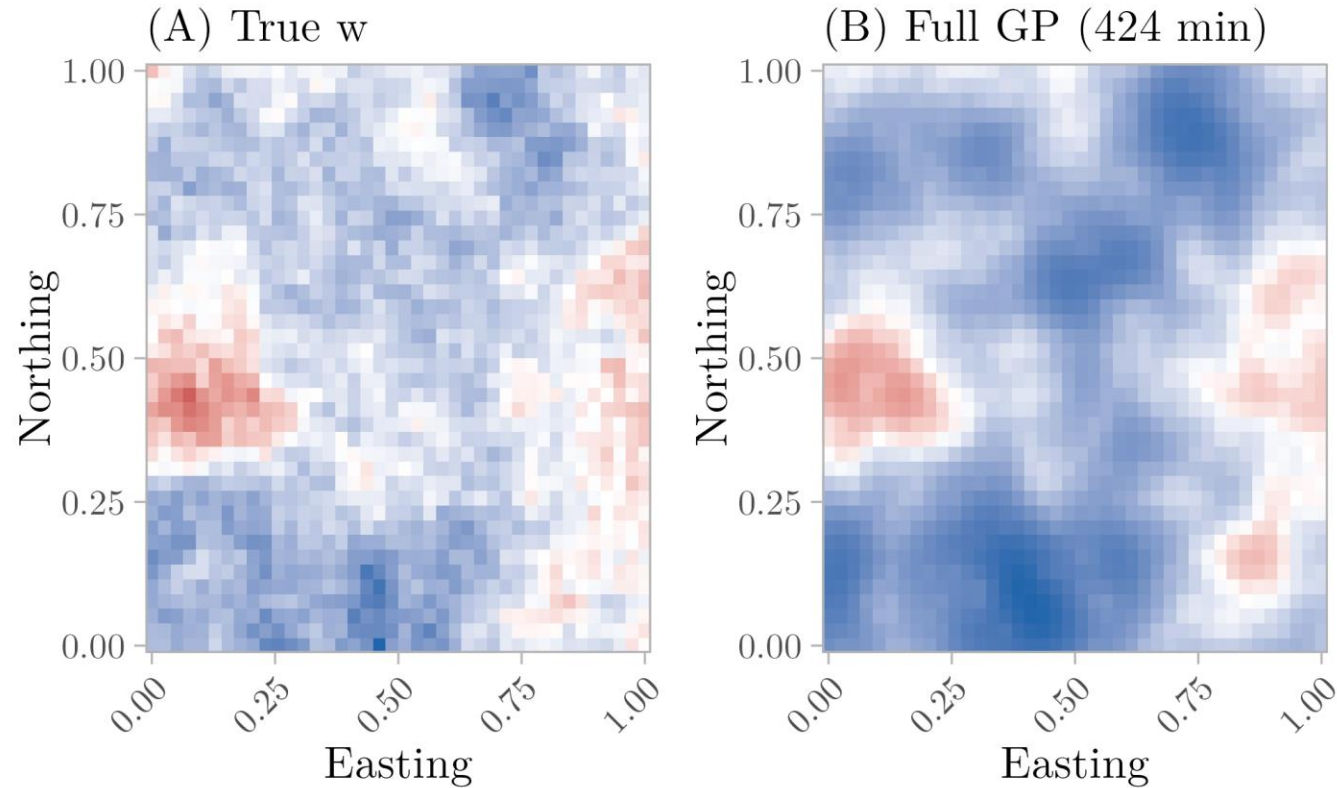
- Estimating the parameters in our model requires calculating:

$$p(\mathbf{w}) \propto -\frac{1}{2}\log(\det(\mathbf{C})) - \frac{1}{2}\mathbf{w}^\top \mathbf{C}^{-1}\mathbf{w}$$

- Storage requirements: $O(J^2)$ flops
- \mathbf{C} is a dense matrix (i.e., there are no zeros)
- Computational complexity: $O(J^3)$ flops
- Computationally infeasible when J is even moderately large (e.g., 500)

Simulation comparison

- $J = 1600$, fit a single-species spatial occupancy model



Solutions to the "big n" problem

A Case Study Competition Among Methods for Analyzing Large Spatial Data

Matthew J. HEATON¹, Abhirup DATTA, Andrew O. FINLEY,
Reinhard FURRER, Joseph GUINNESS, Rajarshi GUHANIYOGI,
Florian GERBER, Robert B. GRAMACY, Dorit HAMMERLING,
Matthias KATZFUSS, Finn LINDGREN, Douglas W. NYCHKA, Furong SUN,
and Andrew ZAMMIT-MANGION

Solutions to the "big n" problem

Low-rank methods:

reduces the dimensionality of the covariance matrix (e.g., analogous to how a PCA takes a set of variables and reduces it to a smaller set)

Solutions to the "big n" problem

Low-rank methods:

reduces the dimensionality of the covariance matrix (e.g., analogous to how a PCA takes a set of variables and reduces it to a smaller set)

Sparse methods:

strategically replace a lot of non-zero values in the covariance (or precision) matrix with 0s.

Solutions to the "big n" problem

Low-rank methods:

reduces the dimensionality of the covariance matrix (e.g., analogous to how a PCA takes a set of variables and reduces it to a smaller set)

Sparse methods:

strategically replace a lot of non-zero values in the covariance (or precision) matrix with 0s.

Algorithmic

methods: focus more on improving algorithms to fit the model as opposed to building the model in a new way.

Solutions to the "big n" problem

Low-rank methods:

reduces the dimensionality of the covariance matrix (e.g., analogous to how a PCA takes a set of variables and reduces it to a smaller set)

Sparse methods:

strategically replace a lot of non-zero values in the covariance (or precision) matrix with 0s.

Algorithmic

methods: focus more on improving algorithms to fit the model as opposed to building the model in a new way.

Nearest Neighbor Gaussian Processes (NNGPs)

- [Datta et al. 2016](#), [Finley et al. 2019](#)

Nearest Neighbor Gaussian Processes (NNGPs)

- [Datta et al. 2016](#), [Finley et al. 2019](#)
- Same parameters and interpretation as the full GP!

Nearest Neighbor Gaussian Processes (NNGPs)

- [Datta et al. 2016](#), [Finley et al. 2019](#)
- Same parameters and interpretation as the full GP!
- Conceptually:
 1. Order the spatial locations (e.g., along the x-axis)

Nearest Neighbor Gaussian Processes (NNGPs)


- [Datta et al. 2016](#), [Finley et al. 2019](#)
- Same parameters and interpretation as the full GP!
- Conceptually:
 1. Order the spatial locations (e.g., along the x-axis)
 2. Determine the m nearest neighbors (subject to ordering) based on Euclidean (linear) distance

Nearest Neighbor Gaussian Processes (NNGPs)

- [Datta et al. 2016](#), [Finley et al. 2019](#)
- Same parameters and interpretation as the full GP!
- Conceptually:
 1. Order the spatial locations (e.g., along the x-axis)
 2. Determine the m nearest neighbors (subject to ordering) based on Euclidean (linear) distance
 3. The spatial random effect at each site only depends on values of its m nearest neighbors and is conditionally independent of all other values


Choosing the neighbors

- spOccupancy and spAbundance order sites along the horizontal axis (i.e., Easting)
- Example: NNGP with 5 neighbors
- Red point denotes the current site
- Blue points denote sites in the "neighbor set"




Choosing the neighbors

- spOccupancy and spAbundance order sites along the horizontal axis (i.e., Easting)
- Example: NNGP with 5 neighbors
- Red point denotes the current site
- Blue points denote sites in the "neighbor set"




Choosing the neighbors

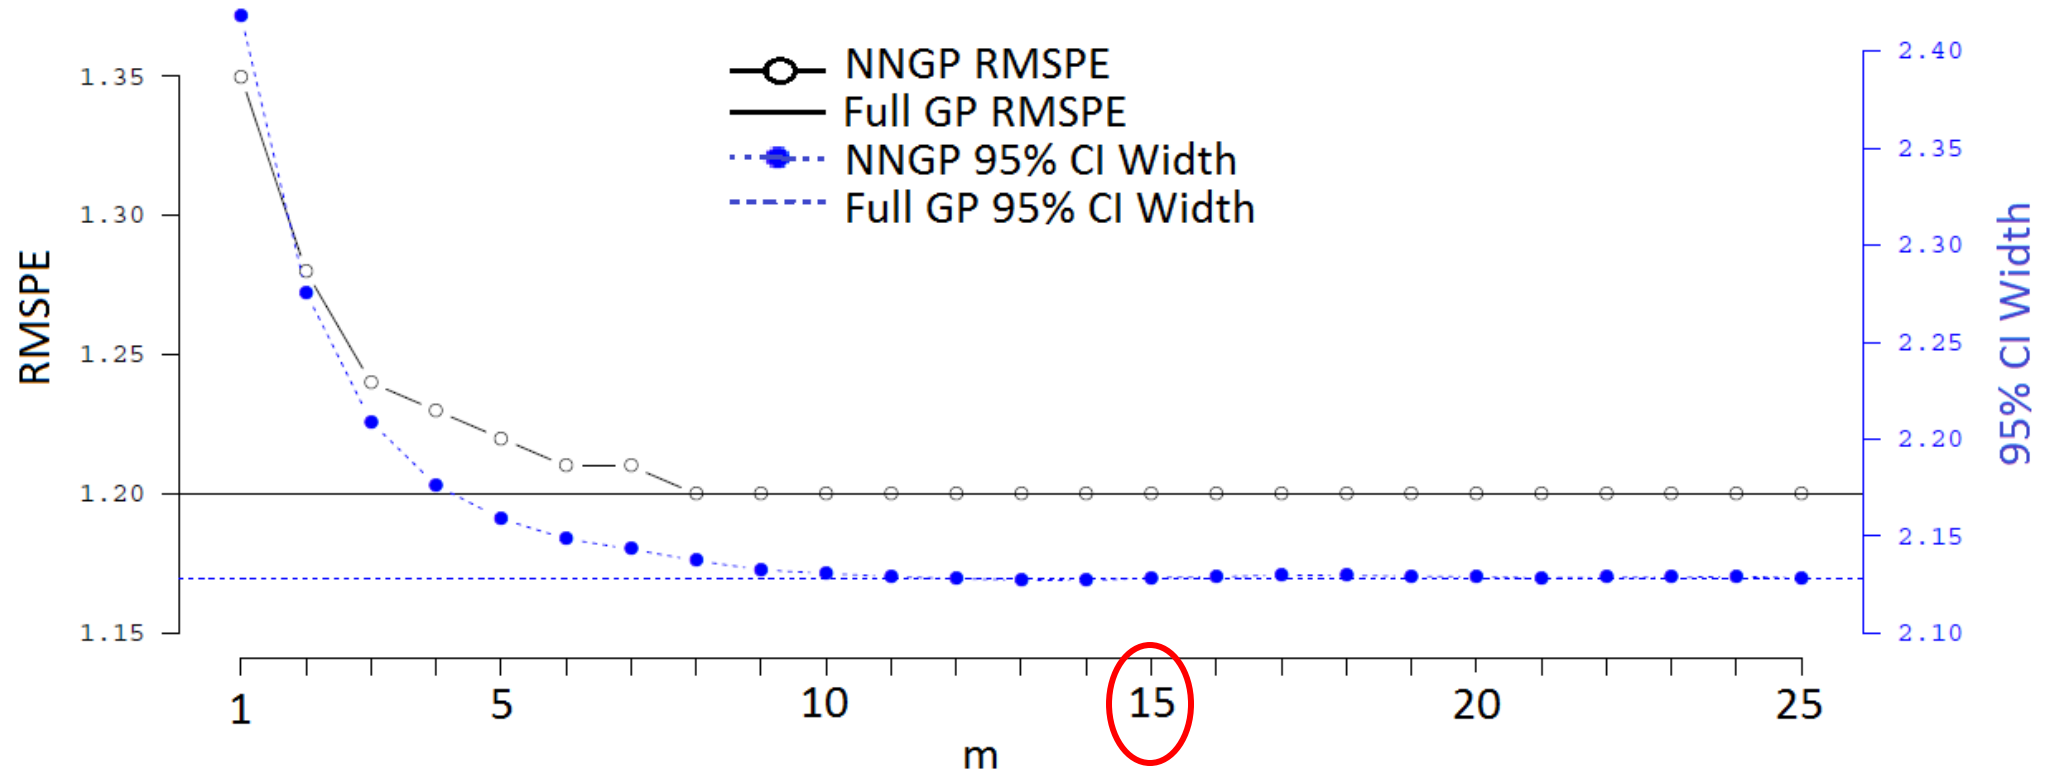
- spOccupancy and spAbundance order sites along the horizontal axis (i.e., Easting)
- Example: NNGP with 5 neighbors
- Red point denotes the current site
- Blue points denote sites in the "neighbor set"



How many neighbors?



How many neighbors?



- $m=15$ neighbors is often adequate (software default)
- Can compare smaller m using WAIC



NNGPs in more detail



NNGPs in more detail


- Use a *sparse* matrix to approximate the dense Gaussian Process covariance matrix \mathbf{C} (Datta et al. 2016)
- Based on rewriting the GP for \mathbf{w} as a product of conditional densities.



NNGPs in more detail

- Use a *sparse* matrix to approximate the dense Gaussian Process covariance matrix \mathbf{C} (Datta et al. 2016)
- Based on rewriting the GP for \mathbf{w} as a product of conditional densities.
- $\mathbf{w} \sim \text{Normal}(\mathbf{0}, \mathbf{C})$ is equivalent to


$$p(w_1)p(w_2 \mid w_1)p(w_3 \mid w_1, w_2) \cdots p(w_n \mid w_1, w_2, \dots, w_{n-1})$$



Introducing sparsity via graphical models




Introducing sparsity via graphical models



$$\begin{aligned} & p(w_1)p(w_2 \mid w_1)p(w_3 \mid w_1, w_2)p(w_4 \mid w_1, w_2, w_3) \\ & \times p(w_5 \mid w_1, w_2, w_3, w_4)p(w_6 \mid w_1, w_2, \dots, w_5) \\ & \times p(w_7 \mid w_1, w_2, \dots, w_6) . \end{aligned}$$




Introducing sparsity via graphical models



$$\begin{aligned} & p(w_1)p(w_2 \mid w_1)p(w_3 \mid w_1, w_2)p(w_4 \mid w_1, w_2, w_3) \\ & \times p(w_5 \mid \cancel{w_1}, w_2, w_3, w_4)p(w_6 \mid w_1, \cancel{w_2}, \cancel{w_3}, w_4, w_5) \\ & \times p(w_7 \mid w_1, w_2, \cancel{w_3}, \cancel{w_4}, \cancel{w_5}, w_6) \end{aligned}$$



Introducing sparsity via graphical models




Size of neighbor set
is $\leq m$

$$\begin{aligned} & p(w_1)p(w_2 \mid w_1)p(w_3 \mid w_1, w_2)p(w_4 \mid w_1, w_2, w_3) \\ & \times p(w_5 \mid \cancel{w_1}, w_2, w_3, w_4)p(w_6 \mid w_1, \cancel{w_2}, \cancel{w_3}, w_4, w_5) \\ & \times p(w_7 \mid w_1, w_2, \cancel{w_3}, \cancel{w_4}, \cancel{w_5}, w_6) \end{aligned}$$



Sparsity via NNGP

$$\text{Normal}(\mathbf{w} \mid \mathbf{0}, \mathbf{C}) \approx \text{Normal}(\mathbf{w} \mid \mathbf{0}, \tilde{\mathbf{C}})$$



NNGP-derived
covariance matrix

Bayesian spatial NNGP linear model

$$y(\mathbf{s}_j) \sim \text{Normal}(\mu(\mathbf{s}_j), \tau^2)$$

$$\mu(\mathbf{s}_j) = \beta_0 + \beta_1 \cdot x_1(\mathbf{s}_j) + \mathbf{w}(\mathbf{s}_j)$$

$$\mathbf{w} \sim \text{Multivariate Normal}(\mathbf{0}, \tilde{\mathbf{C}}(d, \sigma^2, \phi))$$

- All we do is replace the covariance matrix from the GP with the NNGP-derived covariance matrix. Otherwise, it's exactly the same!




Storage and computation

- Storage
 - Never need to store $J \times J$ distance matrix
 - Stores smaller $m \times m$ matrices (m is the number of neighbors)
 - Total storage requirements: $O(Jm^2)$
- Computation:
 - Only involves inverting small $m \times m$ matrices
 - $O(Jm^3)$ flops
- Since $m \ll J$, NNGP offers great scalability for big spatial data

Simulation comparison

- $J = 1600$, fit a single-species spatial occupancy model



Simulation comparison

3a-gp-sim-example.R
3b-nngp-sim-example.R



- $J = 1600$, fit a single-species spatial occupancy model

