



Hierarchical spatial modelling for applied population and community ecology

Jeffrey W. Doser, Marc Kéry,
Gesa von Hirschheydt

24-27 June 2024





Spatial N- Mixture Models

Jeffrey W. Doser
24-27 June 2024



Recap: why N-mixture models?

- An *unmarked* protocol for estimating population abundance.
- Only requires repeat counts of individuals, don't need to uniquely mark individuals.
- Conceptually very similar to occupancy models: variation in observed counts during a period of closure allows for estimation of detection probability.

N-mixture model assumptions

1. **Closure:** all within-site variation in counts is attributable to detection probability.

N-mixture model assumptions

1. **Closure:** all within-site variation in counts is attributable to detection probability.
2. **No false positive errors**

N-mixture model assumptions

1. **Closure:** all within-site variation in counts is attributable to detection probability.
2. **No false positive errors**
3. **Detections are independent**

N-mixture model assumptions

1. **Closure:** all within-site variation in counts is attributable to detection probability.
2. **No false positive errors**
3. **Detections are independent**
4. **No individual variation in detection probability**

N-mixture model assumptions

1. **Closure:** all within-site variation in counts is attributable to detection probability.
2. **No false positive errors**
3. **Detections are independent**
4. **No individual variation in detection probability**
5. **Parametric modelling assumptions:** do the distributions we choose adequately represent the data-generating process?

N-mixture model

$$N_j \sim \text{Poisson}(\mu_j)$$

$$y_{j,k} \sim \text{Binomial}(N_j, p_{j,k})$$

$$\log(\mu_j) = \mathbf{X}_j^\top \boldsymbol{\beta}$$

$$\text{logit}(p_{j,k}) = \mathbf{V}_{j,k}^\top \boldsymbol{\alpha}$$

N-mixture model

$$N_j \sim \text{Poisson}(\mu_j)$$

$$y_{j,k} \sim \text{Binomial}(N_j, p_{j,k})$$

$$\log(\mu_j) = \mathbf{X}_j^\top \boldsymbol{\beta}$$

$$\text{logit}(p_{j,k}) = \mathbf{V}_{j,k}^\top \boldsymbol{\alpha}$$

Matrix notation



$$\log(\mu_j) = \beta_1 + \beta_2 \cdot X_{2,j} + \cdots + \beta_r \cdot X_{r,j}$$

N-mixture model

Can switch to negative binomial

$$N_j \sim \text{Poisson}(\mu_j)$$

$$y_{j,k} \sim \text{Binomial}(N_j, p_{j,k})$$

$$\log(\mu_j) = \mathbf{X}_j^\top \boldsymbol{\beta}$$


$$\text{logit}(p_{j,k}) = \mathbf{V}_{j,k}^\top \boldsymbol{\alpha}$$


Matrix notation

$$\log(\mu_j) = \beta_1 + \beta_2 \cdot X_{2,j} + \cdots + \beta_r \cdot X_{r,j}$$

Fitting N-mixture models in spAbundance

- `NMix()` function
- Family argument can be Poisson or NB (aiming to have zero-inflated Poisson in the future)
- Can include both random intercepts and random slopes (random slopes not yet available in `spOccupancy`). Note that if including both, they are modeled as uncorrelated.
- Will usually need more MCMC iterations to achieve convergence than a comparable occupancy model fit in `spOccupancy`.





Exercise: Estimating tropical bird abundance across an elevational gradient

11a-crimson-mantled-woodpecker-nmix.R



Making N-mixture models spatial

- Our approach for "making N-mixture models spatial" (i.e., accounting for residual spatial autocorrelation) is identical to our spatial occupancy approach.

Making N-mixture models spatial

- Our approach for "making N-mixture models spatial" (i.e., accounting for residual spatial autocorrelation) is identical to our spatial occupancy approach.
- Only focusing on spatial autocorrelation in latent abundance, not in detection.

Making N-mixture models spatial

- Our approach for "making N-mixture models spatial" (i.e., accounting for residual spatial autocorrelation) is identical to our spatial occupancy approach.
- Only focusing on spatial autocorrelation in latent abundance, not in detection.
- Count data provide us with potentially MUCH more information for estimating spatial random effects than detection-nondetection data.

Making N-mixture models spatial

- Our approach for "making N-mixture models spatial" (i.e., accounting for residual spatial autocorrelation) is identical to our spatial occupancy approach.
- Only focusing on spatial autocorrelation in latent abundance, not in detection.
- Count data provide us with potentially MUCH more information for estimating spatial random effects than detection-nondetection data.
- As a result, may often find less uncertainty in spatial random effects in abundance models compared to occupancy models with the same number of spatial locations.

Spatial N-mixture models

$$N_j \sim \text{Poisson}(\mu_j)$$

$$y_{j,k} \sim \text{Binomial}(N_j, p_{j,k})$$

$$\log(\mu_j) = \mathbf{X}_j^\top \boldsymbol{\beta} + w_j$$

$$\mathbf{w} \sim \text{Multivariate Normal}(\mathbf{0}, \tilde{\mathbf{C}}(d, \phi, \sigma^2))$$

$$\text{logit}(p_{j,k}) = \mathbf{V}_{j,k}^\top \boldsymbol{\alpha}$$

Interpretation of spatial covariance parameters

- The additional information in count data generally leads to more precise estimates of the spatial covariance parameters.

Interpretation of spatial covariance parameters

- The additional information in count data generally leads to more precise estimates of the spatial covariance parameters.
- Remember our interpretation of the spatial decay ϕ when using an exponential covariance model

Interpretation of spatial covariance parameters

- The additional information in count data generally leads to more precise estimates of the spatial covariance parameters.
- Remember our interpretation of the spatial decay ϕ when using an exponential covariance model

$$\text{ESR} = \frac{3}{\phi}$$

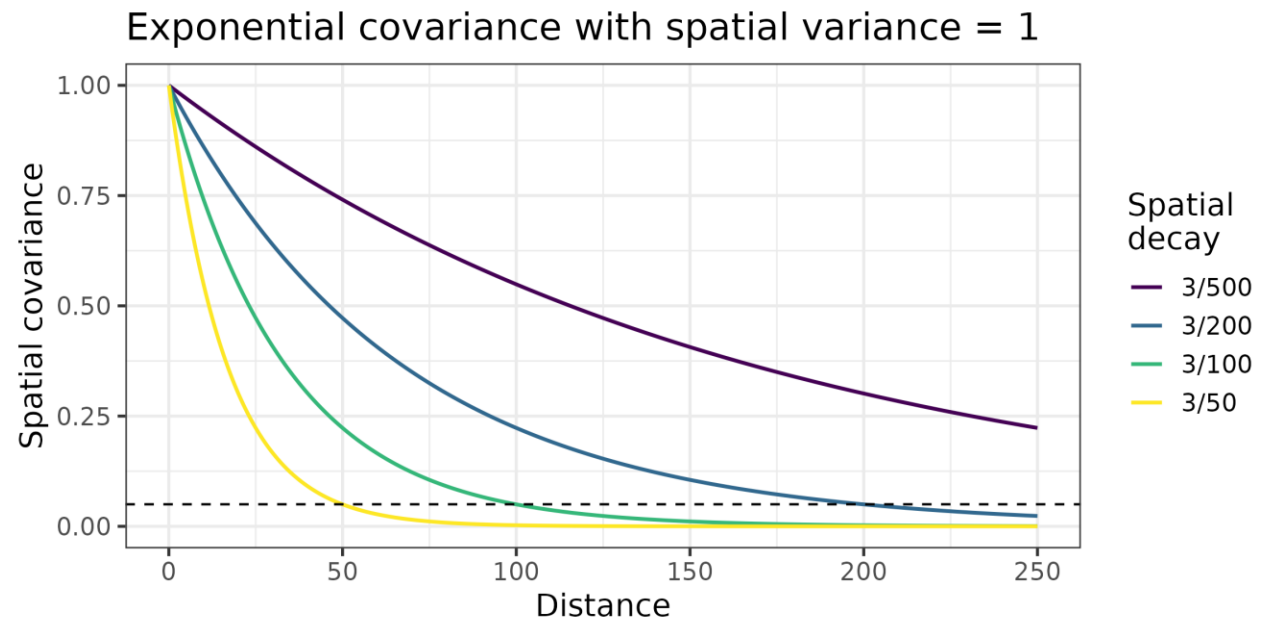
Effective spatial range

Interpretation of spatial covariance parameters

- The additional information in count data generally leads to more precise estimates of the spatial covariance parameters.
- Remember our interpretation of the spatial decay ϕ when using an exponential covariance model

$$\text{ESR} = \frac{3}{\phi}$$

Effective spatial range



Posterior predictive checks for N-mixture models

Posterior predictive checks for N-mixture models

1. Fit the model and generate replicate data values from the posterior predictive distribution.

Posterior predictive checks for N-mixture models

1. Fit the model and generate replicate data values from the posterior predictive distribution.
2. Optionally bin the replicate and actual data values in some way (e.g., sum them across sites or across replicates at each site), or just leave them as is.

Posterior predictive checks for N-mixture models

1. Fit the model and generate replicate data values from the posterior predictive distribution.
2. Optionally bin the replicate and actual data values in some way (e.g., sum them across sites or across replicates at each site), or just leave them as is.
3. Compute a fit statistic on both the actual data and the model-generated replicate data

Posterior predictive checks for N-mixture models

1. Fit the model and generate replicate data values from the posterior predictive distribution.
2. Optionally bin the replicate and actual data values in some way (e.g., sum them across sites or across replicates at each site), or just leave them as is.
3. Compute a fit statistic on both the actual data and the model-generated replicate data
4. Compare the fit statistics for the true and replicate data. If they are widely different, this suggests inadequate model fit.

Posterior predictive checks for N-mixture models

- For N-mixture models, there are two ways we can generate the replicate data values
- In `spAbundance`, we refer to the two approaches as *marginal* replicate data values, or *conditional* replicate data values.
- **Conditional replicate data values:** replicate values are generated conditional on the latent abundance N values.
- **Marginal replicate data values:** replicate values are not generated conditional on the latent abundance N values.

Generating conditional fitted values

$$y_{\text{rep},j,k}^{(l)} \sim \text{Binomial}(N_j^{(l)}, p_{j,k}^{(l)})$$

Generating conditional fitted values

$$\text{MCMC iteration} \longrightarrow y_{\text{rep},j,k}^{(l)} \sim \text{Binomial}(N_j^{(l)}, p_{j,k}^{(l)})$$

Generating conditional fitted values

$$\text{MCMC iteration} \longrightarrow y_{\text{rep},j,k}^{(l)} \sim \text{Binomial}(N_j^{(l)}, p_{j,k}^{(l)})$$

- A replicate data value is generated for each MCMC iteration (l).
- The replicate data are conditional on N_j
- What is the minimum value that $N_j^{(l)}$ can take?

Generating marginal fitted values

$$\hat{N}_{\text{rep},j}^{(l)} \sim \text{Poisson}(\mu_j^{(l)})$$
$$y_{\text{rep},j,k}^{(l)} \sim \text{Binomial}(\hat{N}_j^{(l)}, p_{j,k}^{(l)})$$

- First predict a value of latent abundance at site j for MCMC iteration l.
- Generate the replicate data value using the predicted abundance value
- New predicted latent abundance value not completely dependent on the observed data.

Marginal vs. conditional posterior predictive checks

- See [vignette](#) for small simulation study.
- Marginal PPCs may be more sensitive, but need to do more simulation analyses.
- Lots of opportunities to explore GoF assessments and PPCs in N-mixture models and other types of hierarchical models.

Generating replicate values and PPCs

Posterior predictive checks

```
ppcAbund(object, fit.stat, group, type = 'marginal', ...)
```

Generating replicate (fitted) values

```
# S3 method for spNMix  
fitted(object, type = 'marginal', ...)
```

Model selection with WAIC

- The `waicAbund()` function calculates WAIC for all `spAbundance` models.

Model selection with WAIC

- The `waicAbund()` function calculates WAIC for all `spAbundance` models.

$$\text{WAIC} = -2 \times (\text{elppd} - \text{pD})$$

Model selection with WAIC

- The `waicAbund()` function calculates WAIC for all `spAbundance` models.

$$\text{WAIC} = -2 \times (\text{elppd} - \text{pD})$$



Expected log
pointwise
predictive density



Effective number
of parameters

Model selection with WAIC

- The `waicAbund()` function calculates WAIC for all `spAbundance` models.

$$\text{WAIC} = -2 \times (\text{elppd} - \text{pD})$$

Expected log
pointwise
predictive density



Effective number
of parameters

- Calculating `elppd` in N-mixture and hierarchical distance sampling models requires integrating out latent abundance
- Calculation of WAIC can be slow, particularly with large counts.

Exercise: Predicting abundance across a simulated landscape

11b-sim-spatial-nmix.R



Difficulties with spatial N-mixture models

- When model assumptions are met and overdispersion in counts is reasonable, N-mixture models can work well and give reasonable estimates of abundance (our analysis of the simulated data).

Difficulties with spatial N-mixture models

- When model assumptions are met and overdispersion in counts is reasonable, N-mixture models can work well and give reasonable estimates of abundance (our analysis of the simulated data).
- When there is large overdispersion, spatial N-mixture models (or NB N-mixture models) can be VERY difficult to fit.

Difficulties with spatial N-mixture models

- When model assumptions are met and overdispersion in counts is reasonable, N-mixture models can work well and give reasonable estimates of abundance (our analysis of the simulated data).
- When there is large overdispersion, spatial N-mixture models (or NB N-mixture models) can be VERY difficult to fit.
- Out of all the models in `spOccupancy`/`spAbundance`, spatial N-mixture models (single-species and multi-species) are the trickiest.

Difficulties with spatial N-mixture models

- When model assumptions are met and overdispersion in counts is reasonable, N-mixture models can work well and give reasonable estimates of abundance (our analysis of the simulated data).
- When there is large overdispersion, spatial N-mixture models (or NB N-mixture models) can be VERY difficult to fit.
- Out of all the models in `spOccupancy`/`spAbundance`, spatial N-mixture models (single-species and multi-species) are the trickiest.
- Why?

Difficulties with spatial N-mixture models

- Recall a key goal of N-mixture models: partition variability in detection probability from variability in abundance.
- N-mixture models are most commonly fit with a Binomial detection sub-model and a Poisson abundance sub-model.
- This works well without any overdispersion.


Difficulties with spatial N-mixture models

- Knape et al. (2018) showed that an N-mixture model with overdispersion in detection (beta-binomial Poisson N-mixture) and a model with overdispersion in abundance (binomial negative binomial N-mixture) are nearly identical.




Received: 23 January 2018 | Accepted: 2 July 2018

DOI: 10.1111/2041-210X.13062

RESEARCH ARTICLE

Methods in Ecology and Evolution  BRITISH
ECOLOGICAL
SOCIETY

Sensitivity of binomial N-mixture models to overdispersion: The importance of assessing model fit

Jonas Knape¹  | Debora Arlt¹ | Frédéric Barraquand²  | Åke Berg¹ | Mathieu Chevalier¹ | Tomas Pärt¹ | Alejandro Ruete^{1,3}  | Michał Żmihorski¹

Difficulties with spatial N-mixture models

- Knape et al. (2018) showed that an N-mixture model with overdispersion in detection (beta-binomial Poisson N-mixture) and a model with overdispersion in abundance (binomial negative binomial N-mixture) are nearly identical.




Received: 23 January 2018 | Accepted: 2 July 2018

DOI: 10.1111/2041-210X.13062

RESEARCH ARTICLE

Methods in Ecology and Evolution  BRITISH
ECOLOGICAL
SOCIETY

Sensitivity of binomial N-mixture models to overdispersion: The importance of assessing model fit

Jonas Knape¹  | Debora Arlt¹ | Frédéric Barraquand²  | Åke Berg¹ | Mathieu Chevalier¹ | Tomas Pärt¹ | Alejandro Ruete^{1,3}  | Michał Żmihorski¹

It is very difficult to distinguish between overdispersion in detection probability and overdispersion in abundance.

Difficulties with spatial N-mixture models

- As a result, spatial N-mixture models can be extremely slow to converge and may require heavily informative prior distributions.
- Further, overdispersed models (i.e., NB N-mixture models and spatial N-mixture models) can yield very high and unreasonable estimates of abundance
 - Often happens in situations where the closure assumption may be violated
 - Kéry (2010) Ecology
- **What should we do with heavily overdispersed repeated count data?**


Solutions

1. Explore the use of more informative prior distributions on spatial parameters or the NB dispersion parameter.
2. Consider using GLMMs and estimating relative abundance.

Further complications with N-mixture models


- N-mixture models are powerful, but very sensitive to model assumptions
 - "With great power comes great responsibility."
- Closure violations can lead to biased abundance estimates (i.e., estimates will be biased high).
- Lots of discussion: Duarte et al. (2018), Knape et al. (2018), Bellier et al. (2016), Link et al. (2018), Barker et al. (2018), and many others.
- In cases where assumptions may be violated, abundance estimates from N-mixture models may best be viewed as relative abundance.

N-mixture models vs. GLMMs for relative abundance estimation




Goldstein and de Valpine (2022) *Sci Reports*

N-mixture models vs. GLMMs for relative abundance estimation



Lots of outstanding questions to understand performance of N-mixture models for relative abundance estimation

Goldstein and de Valpine (2022) *Sci Reports*



Exercise: Exploring difficulties with spatial N-mixture models

11c-european-goldfinch-spatial-nmix.R

